

Single Machine:**16 Cores:****Output:**

Parameters: 128, 4, 4, 4, 15

Worker 1: 4 matrix multiplies performed, 0.07000000 time working in seconds

Worker 2: 5 matrix multiplies performed, 0.06000000 time working in seconds

Total time: 0.07000000

Parameters: 64, 8, 8, 8, 15

Worker 1: 34 matrix multiplies performed, 0.05000000 time working in seconds

Worker 2: 34 matrix multiplies performed, 0.05000000 time working in seconds

Total time: 0.05000000

Parameters: 32, 16, 16, 16, 15

Worker 1: 273 matrix multiplies performed, 0.31000000 time working in seconds

Worker 2: 273 matrix multiplies performed, 0.30000000 time working in seconds

Total time: 0.33000000

(I omitted the output for processes 3-15 to save on space but had similar output)

MPI command: mpirun -mca btl self,sm,tcp -mca btl_tcp_if_include 132.249.107.0/24 -map-by node -np 16 mmmw 32 16 16 16 mmmw-1-15.out >> mmmw-1-15-data.out

Network: 10Gbe

Expectations: I expect that with a large enough block size, 15 workers will perform the best but if the block size gets too small then 15 workers will perform worse than the original.

Speedup/Slowdown: $0.3447 / 0.07 = 4.92$, $0.3447 / 0.05 = 6.894$, $0.3447 / 0.33 = 1.04$

Efficiency: $4.92 / 16 = 0.3075$, $6.894 / 16 = 0.43$, $1.04 / 16 = 0.065$

Conclusion: My expectations were correct in that it looks like the sweet spot for 15 workers is around a block size of 64 bytes.

9 Cores:**Output:**

Parameters: 128, 4, 4, 4, 8

Worker 1: 8 matrix multiplies performed,	0.10000000 time working in seconds
Worker 2: 8 matrix multiplies performed,	0.13000000 time working in seconds
Worker 3: 9 matrix multiplies performed,	0.11000000 time working in seconds
Worker 4: 8 matrix multiplies performed,	0.22000000 time working in seconds
Worker 5: 8 matrix multiplies performed,	0.22000000 time working in seconds
Worker 6: 7 matrix multiplies performed,	0.12000000 time working in seconds
Worker 7: 8 matrix multiplies performed,	0.23000000 time working in seconds
Worker 8: 8 matrix multiplies performed,	0.22000000 time working in seconds
Total time:	0.22000000

MPI command: mpirun -mca btl self,sm,tcp -mca btl_tcp_if_include 132.249.107.0/24 -map-by node -np 9 mmmw 32 16 16 16 mmmw-1-8.out >> mmmw-1-8-data.out

Network: 10Gbe

Expectations: I expect that with a medium enough block size, 8 workers will perform the best but if the block size gets too small then 8 workers will perform worse than the original.

Speedup/Slowdown: $0.3447 / 0.22 = 1.56$

Efficiency: $1.56 / 9 = 0.17$

Conclusion: I was only able to run 1 test for 8 workers and a large block size increased performance

5 Cores:**Output:**

Parameters: 128, 4, 4, 4, 4

Worker 1: 15 matrix multiplies performed,	0.09000000 time working in seconds
Worker 2: 17 matrix multiplies performed,	0.18000000 time working in seconds
Worker 3: 16 matrix multiplies performed,	0.17000000 time working in seconds
Worker 4: 16 matrix multiplies performed,	0.10000000 time working in seconds
Total time:	0.18000000

Parameters: 64, 8, 8, 8, 4

Worker 1: 129 matrix multiplies performed,	0.04000000 time working in seconds
Worker 2: 127 matrix multiplies performed,	0.04000000 time working in seconds
Worker 3: 129 matrix multiplies performed,	0.05000000 time working in seconds
Worker 4: 127 matrix multiplies performed,	0.04000000 time working in seconds
Total time:	0.05000000

Parameters: 32, 16, 16, 16, 4

Worker 1: 1025 matrix multiplies performed,	0.30000000 time working in seconds
Worker 2: 1025 matrix multiplies performed,	0.29000000 time working in seconds
Worker 3: 1023 matrix multiplies performed,	0.29000000 time working in seconds
Worker 4: 1023 matrix multiplies performed,	0.29000000 time working in seconds
Total time:	0.29000000

MPI command: mpirun -mca btl self,sm,tcp -mca btl_tcp_if_include 132.249.107.0/24 -map-by node -np 5 mmmw 32 16 16 16 mmmw-1-4.out >> mmmw-1-4-data.out

Network: 10Gbe

Expectations: I expect that again a medium block size will perform the best, and large block sizes will be terrible because there are less workers available

Speedup/Slowdown: $0.3447 / 0.18 = 1.915$, $0.3447 / 0.05 = 6.894$, $0.3447 / 0.29 = 1.19$

Efficiency: $1.915 / 5 = 0.383$, $6.894 / 5 = 1.3788$, $1.19 / 5 = 0.238$

Conclusion: My expectations were somewhat correct in that a medium block size performed the best but also a small block size performed the worst.

I only was able to run on a single machine because my jobs weren't running on tsc for multiple nodes for some reason. Also I ran out of time for my report.